

Operating Systems Assignment 2 - Easy

Aditya Jha 2022CS11102
Rishit Jakharia 2022CS11621

1 Signal Handling in xv6

1.1 Console input

Updated the *console.c* file to accept inputs;
Using definitions like `#define SIGINT 1`
Code being as follows;

```
// console.c
void consoleintr(int (*getc)(void)){
...
    switch(c){
        ...
        case C('C'):
            dosignal = SIGINT;
            break;
        case C('B'):
            dosignal = SIGBG;
            break;
        ... // Similarly for Ctrl+F and Ctrl+G
    }
...
    switch(dosignal){
        case SIGCUSTOM:
            cprintf("Ctrl-G is detected by xv6\n");
            customsig();
            break;
        case SIGINT:
            cprintf("Ctrl-C is detected by xv6\n");
            procsignal(SIGINT);
            break;
        ... // Similarly for SIGBG and SIGFG
        default:
            break;
    }
}
```

1.2 SIGINT, SIGBG, and SIGFG

The process structure (`struct proc`) was extended with:
`suspended` and `ever_suspended` flags

Console Handling (*console.c*)

```
case C('C'): dosignal = SIGINT;  
case C('B'): dosignal = SIGBG;  
case C('F'): dosignal = SIGFG;
```

Later processed through:

```
procsignal(SIGINT/SIGBG/SIGFG/...);
```

Signal Specific details

1) SIGINT

Key Effects:

- Terminates all non-system processes
- Cleans up suspended processes via `cleanup_suspended()`
- Immediate process table sweep with priority given to active processes

```
// proc.c void procsignal(int)  
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){  
    if(p->state == UNUSED || p->pid <= 2) continue;  
    p->killed = 1;  
    if(p->state == SLEEPING) p->state = RUNNABLE;  
}  
cleanup_suspended();
```

2) SIGBG

Behavior:

- Suspends non-critical processes
- Maintains process state for later resumption
- Maintains another process state for eventually reaping
- Suspended processes return a special code '-2' in the `wait()` for the shell to run.

Below is the code for setting the process as suspended.

```

// proc.c void procsignal(int)
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    // Skip unused processes and kernel/system processes (pid
    // <= 2).
    if(p->state == UNUSED || p->pid <= 2)
        continue;

    // Skip processes with reserved names.
    if(strcmp(p->name, "sh") == 0 ||
        strcmp(p->name, "init") == 0 ||
        strcmp(p->name, "console") == 0)
        continue;

    // Suspend the process.
    p->suspended = 1;
    p->ever_suspended = 1;

    // Wakeup its parent process, if necessary.
    wakeup1(p->parent);
}

```

The following is how this process is handled in the wait function.

```

// proc.c int wait(void)
...
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->parent != curproc || p->suspended)
        continue;
    havekids = 1;
    if(p->state == ZOMBIE){
        ...
        if(p->ever_suspended)
            return -2;

        return pid;
    }
}
...
if(all_done && havekids){
    // Return special code to indicate suspended children
    // exist
    release(&ptable.lock);
    return -2;
}
...
}
}

```

3) SIGFG

Foreground restoration

```
// proc.c void procsignal(int)
case SIGFG:
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state != UNUSED && p->suspended){
            p->suspended = 0;
            if(p->state == SLEEPING && !p->killed){
                p->state = RUNNABLE;
            }
        }
    }
}
```

1.3 SIGCUSTOM

Registering the signal handler

For registering a signal handler, a system call is added as follows:

```
// syscall.h
#define SYS_signal 22

// syscall.c
extern int sys_signal(void);

... // (*syscalls)
[SYS_signal] sys_signal,
...

// usys.S
SYSCALL(signal)
```

To keep track of the signal handler for each process, added the handler to the process structure. Then defined a user-space wrapper function.

```
// proc.h
typedef void signal(sighandler_t handler);

struct proc{ ...
    sighandler_t handler;
    ... // some other useful fields
}

// user.h
void signal(sighandler_t handler);
```

When the syscall is made, the address of the handler is saved to the handler attribute of the process.

```
// sysproc.c
int sys_signal(void){
    int addr;
    if(argint(0, &addr) < 0)
        return -1;

    myproc()->handler = (sighandler_t)addr;
    return 0;
}
```

When the console registers a Ctrl+G interrupt, it runs the function void customsig(void) in *proc.c*.

```
// proc.c void customsig(void)
if(p->handler){
    // Save current user context so that sigret can restore it.
    p->saved_eip = p->tf->eip;
    p->saved_esp = p->tf->esp;
    p->in_handler = 1;

    // Adjust the user stack: make room for one return address.
    uint *ustack = (uint*)p->tf->esp;
    ustack -= 1; // Reserve space for return address.
    ustack[0] = 0;

    p->tf->esp = (uint)ustack;
    p->tf->eip = (uint)p->handler;
}
```

The `p->saved_eip` stores the return address and similarly for the stack pointer(`p->saved_esp`). The `p->in_handler` flag is used for the return from the handler.

```
// trap.c void trap(struct trapframe *tf)

default:
    // When the custom handler completes execution,
    // the pointers are restored for the trapframe
    tf->eip = p->saved_eip;
    tf->esp = p->saved_esp;
    p->in_handler = 0;
    break;
```

2 xv6 Scheduler

2.1 Implementation for the syscalls

The `custom_fork` system call extends the standard fork operation with two parameters:

- `start_later_flag`: Controls whether the process starts immediately (0) or is created but not yet runnable (1)
- `exec_time`: Limits the maximum execution time of the process (-1 for unlimited execution)

Key implementation details include:

- Adding a new process state `CREATED` for processes that are fully initialized but not yet runnable
- Implementing a `check_exec_time` function called from the timer interrupt handler to enforce execution time limits
- Handling the special case where `start_later_flag = 0` and `exec_time = -1` by directly calling the standard fork

The `scheduler_start` system call allows processes in the `CREATED` state to become `RUNNABLE`.

```
int sys_scheduler_start(void){
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state == CREATED)
            p->state = RUNNABLE;
    }
    release(&ptable.lock);
    return 0;
}
```

2.2 Priority Boosting Scheduler

The priority boosting scheduler implements a dynamic priority system, calculating process priorities using:

$$\pi_i(t) = \pi_i(0) - \alpha \cdot C_i(t) + \beta \cdot W_i(t)$$

Key implementation details include:

- Adding priority-related fields to the process control block
- Updating process priorities in both the scheduler and the timer interrupt handler
- Implementing highest-priority-first selection with PID as a tiebreaker

To update a process `p`'s priority, the following is done:

```
p->priority = p->init_priority - (ALPHA * p->cpu_ticks) + (BETA *
    p->wait_time);
```

2.3 Effects of α and β Parameters

1) Effects on Profiled Parameters

- Turnaround Time (TAT): Higher α/β ratio generally reduces average TAT by preventing excessive waiting.
- Waiting Time (WT): Higher β values directly reduce average WT.
- Response Time (RT): Higher β values improve RT by giving waiting processes a priority boost.
- Context Switches (#CS): Higher α and β values generally increase #CS.

2) Effects on CPU-bound Jobs

- Higher α values significantly reduce the priority of CPU-bound processes as they accumulate CPU time. This prevents CPU-bound jobs from monopolizing the processor.
- As α increases, CPU-bound jobs experience increased turnaround times, more frequent preemption, longer overall completion times and higher waiting times.
- CPU-bound jobs benefit less from β since they rarely wait voluntarily. This prevents complete starvation of CPU-intensive tasks.

3) Effects on I/O-bound Jobs

- I/O-bound processes accumulate less CPU time, so they're less affected by α .
- These processes benefit significantly from β as they frequently wait.
- As β increases, these jobs experience lower response times and smaller waiting-to-execution time ratios.

The current parameters ($\alpha = 1$, $\beta = 2$) create a scheduler that provides good interactive performance while still allowing background CPU-intensive tasks to make reasonable progress.

Note: Code snippets in the report might differ from what is actually implemented in the xv6 kernel. The snippets here are kept simpler for easier understanding.