

# COL362/632 Introduction to Database Management Systems

## Indexing – B+Trees


Kaustubh Beedkar

Department of Computer Science and Engineering  
Indian Institute of Technology Delhi




# B/B+Tree Trivia

- ▶ Rudolf Bayer (invented B-trees in 1972; with McCreight)
- ▶ Gold standard for indexing,  
a.k.a. the greatest data structure of all times!
- ▶ Impact
  - Databases
  - FileSystems
  - Key Value stores
  - Search Engines
- ▶ No official explanation for what **B** stands for in B-Trees



“the more you think about what the B in B-Tree means, the better you understand B-Trees!”


# Binary Search Trees



Node Structure

P	80	P
---	----	---

# *m*-Way Search Trees




Node Structure






# B-Trees

- ▶  $m$ -way search trees
- ▶ Root node has at least two children (unless it a leaf)
- ▶ Every inner node has at least  $\lceil \frac{m}{2} \rceil$  children
- ▶ All leaf nodes have the same number of ancestors



# B-Trees as a Database Index



# B-Tree Family

**Generally referred to as class of balanced tree data structures**

- ▶ B-Tree
- ▶ B+Tree
- ▶ B\*Tree
- ▶ B<sup>link</sup>Tree
- ▶ B $\epsilon$ -Trees
- ▶ ...

# B+Tree


- ▶ Self balancing  $m$ -way search tree
- ▶ Allows searches, insertions, deletion, and sequential access in  $O(\log n)$
- ▶ Optimized for reading and writing large blocks of data

## Properties

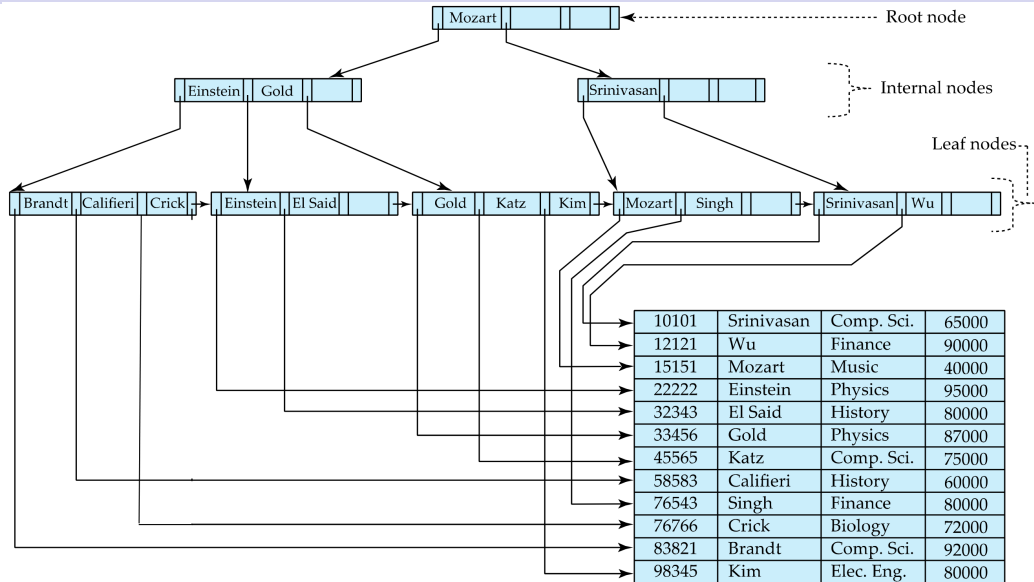
- ▶ All leaf nodes have the same number of ancestors
- ▶ Root node has at least 2 and at most  $m$  children
- ▶ Every leaf node has  $\left\lceil \frac{m-1}{2} \right\rceil \leq \#keys \leq m-1$
- ▶ Every internal node has  $\left\lfloor \frac{m-1}{2} \right\rfloor \leq \#keys \leq m-1$



# B+Tree Example



# B+ Tree (Example from Book B1)



# Node Structure


## Non-leaf Nodes

- ▶ Node with  $K$  keys has  $K + 1$  children



## Leaf Nodes

- ▶ Node with  $K$  keys has  $K$  **record pointers**
- ▶ Has pointers to previous and next leaf notes




# B+Tree Operations

- ▶ Search: point and range queries
- ▶ Insertions
- ▶ Deletions
- ▶ Bulk loading

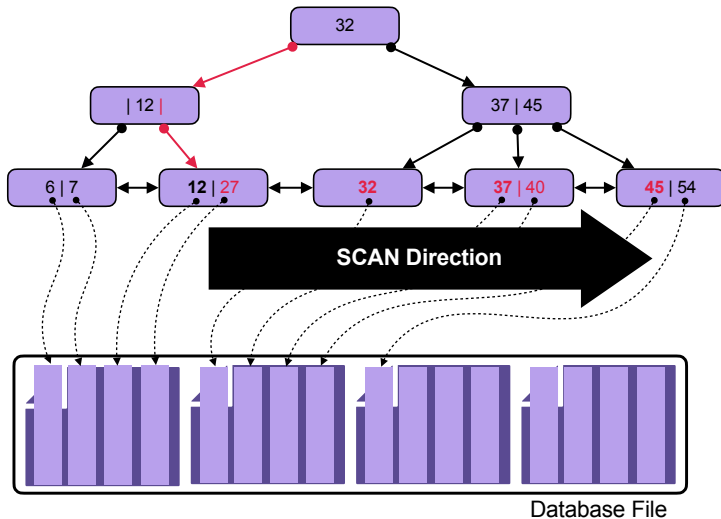
# Point Queries

key=40



# Range Queries


$27 \leq \text{key} \leq 45$



# Clustered and non-clustered B+Tree

## Clustered B+Tree


- Find the left most entry and sequentially scan right



# Clustered and non-clustered B+Tree

## Non-clustered B+Tree

- ▶ First find all pages
- ▶ then, sort by page id






# Searching within a node


## 1. Linear Search

e.g., key=20



## 2. Binary Search


e.g., key=13



## 3. Interpolation Search

Improvement over binary search where values are uniformly distributed

e.g., key=18; offset:  $\frac{(18-10) \times 6}{20-10}$



# B+Tree Insertion

1. Perform search to find the leaf node  $\ell$
2. Check if  $\ell$  is full or not
  - If  $\ell$  is not full, then add the record
  - Otherwise split  $\ell$  into
    - $\ell$  with  $\lfloor \frac{K+1}{2} \rfloor$
    - $\ell'$  with  $\lceil \frac{K+1}{2} \rceil$
3. Insert  $\lceil \frac{K+1}{2} \rceil$ -th key to parent pointing to  $\ell'$
4. Propagate recursively


# Examples

<https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

# Bulk Insert

**Given a collection of records, create a B+tree index on some key**

- ▶ Point inserts are expensive!
- ▶ Instead perform bulk loading
  1. Sort entries according to key
  2. Construct tree bottom-up
- ▶ **Example**  $m = 4$ 
  - Bulk insert: 27, 12, 7, 6, 32, 37
  - Sorted keys: 6, 7, 12, 27, 32, 37



# B+Tree Deletions

1. Perform search to find the leaf node  $\ell$
2. If  $\ell$  is at least half full, delete the entry, done!
3. Otherwise, try redistributing by borrowing from sibling
  - If successful, update the parent node
  - If not, merge<sup>1</sup>  $\ell$  and sibling
    - update parent by deleting the index key pointing to deleted entry

---

<sup>1</sup>Some DBMS may delay merging

# Examples

<https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>

## Handling duplicates

1. Append Record ids to keys
2. Overflow pages

## Handling nulls

- Store all null keys at either first or last leaf nodes

## Handling variable length keys

1. Use pointers (useful for in-memory DBMS)
2. Variable length nodes (requires intricate memory management)
3. Padding (make all keys of the same length)

