

COL362/632 Introduction to Database Management Systems

Indexing – Hash Tables, Bitmaps, & Bloom Filters

Kaustubh Beedkar

Department of Computer Science and Engineering
Indian Institute of Technology Delhi



Outline

1 Hash Index

2 Bitmap and Join Indexes

3 Bloom Filters

Hash Index

- ▶ Widely used for indexes in main memory
- ▶ but, hash file organization is not very widely used
- ▶ Given: set K of keys and B buckets
- ▶ hash function $h : K \mapsto B$
- ▶ Insert: key k , compute $h(k)$ and add entry to the offset
- ▶ Use hash function such that
 - Distribution is uniform
 - Distribution is random
- ▶ Types of hash indexing
 1. Static Hashing
 2. Dynamic Hashing

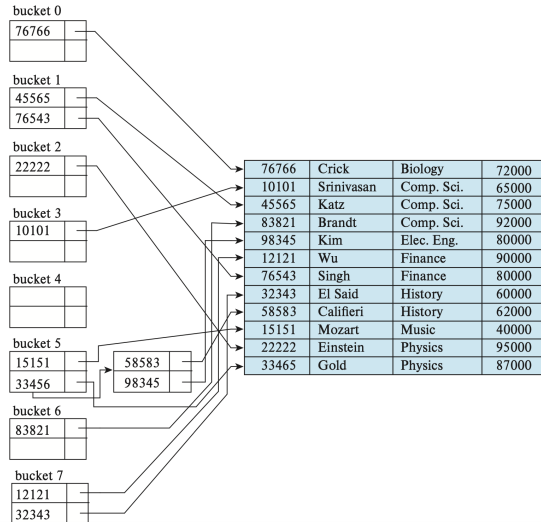
Hashing Index

- ▶ Records are stored in **buckets**
 - linked list of disk blocks
 - linked list of index entries or records (for in-memory indexes)
- ▶ Hash file organization
 - Store actual records instead of record pointers

Bucket Overflows

- ▶ Occur due to
 - Insufficient buckets
 - Data skew or partitioning skew
- ▶ Use **overflow chaining** to handle multiple records in the same buckets

Example Hash Index



Example Hash Organization

bucket 0

| | | | |
|--|--|--|--|
| | | | |
| | | | |
| | | | |
| | | | |

bucket 1

| | | | |
|-------|--------|-------|-------|
| 15151 | Mozart | Music | 40000 |
| | | | |
| | | | |
| | | | |

bucket 2

| | | | |
|-------|-----------|---------|-------|
| 32343 | El Said | History | 80000 |
| 58583 | Califieri | History | 60000 |
| | | | |
| | | | |

bucket 3

| | | | |
|-------|----------|------------|-------|
| 22222 | Einstein | Physics | 95000 |
| 33456 | Gold | Physics | 87000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| | | | |

bucket 4

| | | | |
|-------|-------|---------|-------|
| 12121 | Wu | Finance | 90000 |
| 76543 | Singh | Finance | 80000 |
| | | | |
| | | | |

bucket 5

| | | | |
|-------|-------|---------|-------|
| 76766 | Crick | Biology | 72000 |
| | | | |
| | | | |
| | | | |

bucket 6

| | | | |
|-------|------------|------------|-------|
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| | | | |

bucket 7

| | | | |
|--|--|--|--|
| | | | |
| | | | |
| | | | |
| | | | |

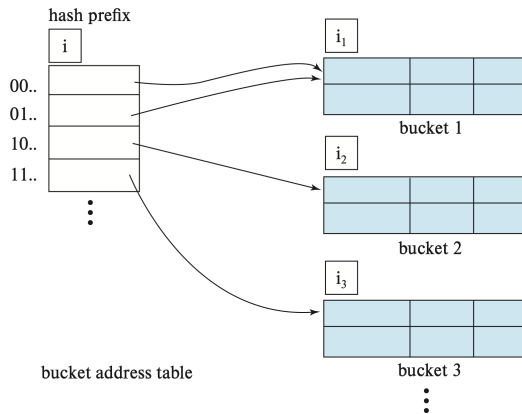
Static Hashing

- ▶ Set of buckets is fixed at the time of index creation (problem!)
 - Long overflow chains as database grows
 - Space wastage as database shrinks
 - Expensive reorganization

Extendible Hashing


- ▶ Type of **dynamic hashing**
- ▶ Keep directory of pointers to buckets
- ▶ Reorganize the index by doubling the directory

Hash Data structure



- ▶ **bucket address table**: directory of pointers to buckets
- ▶ $hash(key) = b$ -bit binary integer (typically $b = 32$)
- ▶ use i higher order bits as an offset
- ▶ use i bits of $hash(key)$ to determine bucket
- ▶ each bucket associated with i_j as length of common hash prefix


Index Construction



| <i>dept_name</i> | <i>h(dept_name)</i> |
|------------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction


1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)



| <i>dept_name</i> | <i>h(dept_name)</i> |
|------------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction


1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)



| <i>dept_name</i> | <i>h(dept_name)</i> |
|------------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction

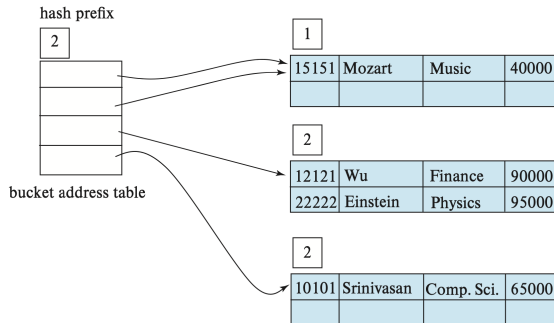
1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)



| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction

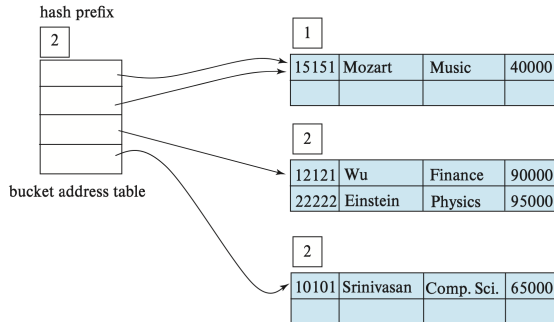
1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)



| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction

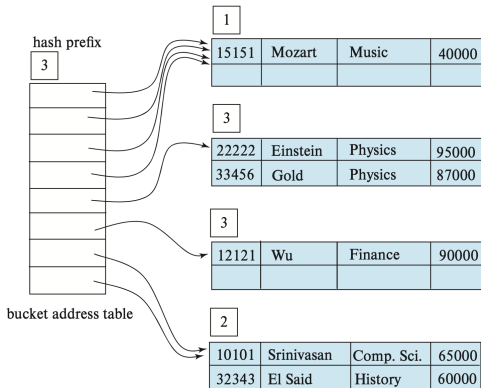
1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)
5. (32343, El Said, History, 60000)
6. (33456, Gold, Physics, 87000)



| <i>dept_name</i> | <i>h(dept_name)</i> |
|------------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction


1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)
5. (32343, El Said, History, 60000)
6. (33456, Gold, Physics, 87000)



| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction

1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)
5. (32343, El Said, History, 60000)
6. (33456, Gold, Physics, 87000)
7. (45565, Katz, Comp. Sci., 75000)

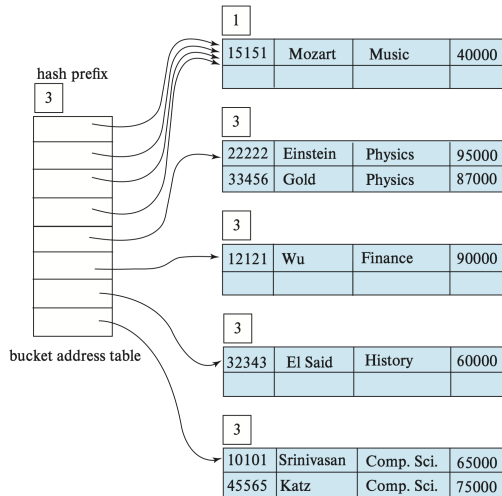


| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Index Construction

1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)
5. (32343, El Said, History, 60000)
6. (33456, Gold, Physics, 87000)
7. (45565, Katz, Comp. Sci., 75000)

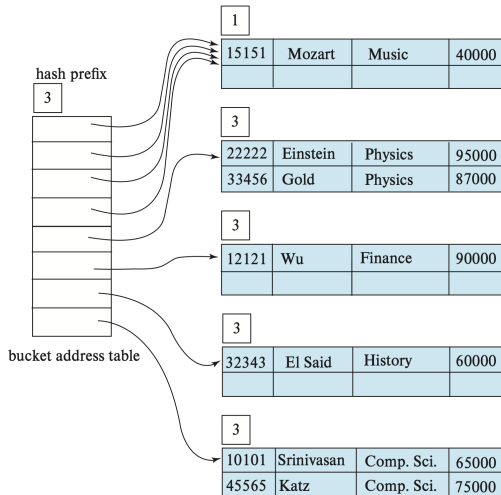
| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |



Index Construction

1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)
5. (32343, El Said, History, 60000)
6. (33456, Gold, Physics, 87000)
7. (45565, Katz, Comp. Sci., 75000)
8. (58583, Califieri, History, 62000)
9. (76543, Singh, Finance, 80000)
10. (76766, Crick, Biology, 72000)
11. (83821, Brandt, Comp. Sci., 92000)

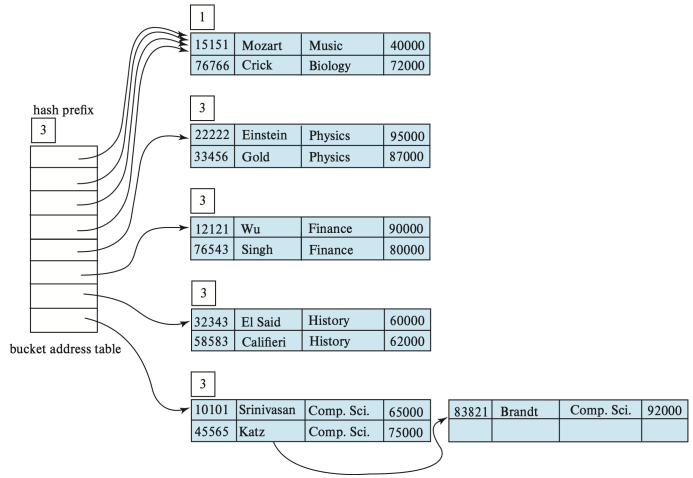
| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |



Index Construction

1. (10101, Srinivasan, Comp. Sci., 65000)
2. (12121, Wu, Finance, 90000)
3. (15151, Mozart, Music, 40000)
4. (22222, Einstein, Physics, 95000)
5. (32343, El Said, History, 60000)
6. (33456, Gold, Physics, 87000)
7. (45565, Katz, Comp. Sci., 75000)
8. (58583, Califieri, History, 62000)
9. (76543, Singh, Finance, 80000)
10. (76766, Crick, Biology, 72000)
11. (83821, Brandt, Comp. Sci., 92000)

| dept_name | h(dept_name) |
|------------|-----------------------------------------|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| Comp. Sci. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| Elec. Eng. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |



Outline

1 Hash Index

2 Bitmap and Join Indexes

3 Bloom Filters

Bitmap Index

- ▶ A **bitmap index** for a field F is a collection of bit-vectors of length n , one for each possible value that may appear in the field F

- ▶ Example: Book table:

| RowID | BookID | Title | Binding | Language |
|-------|--------|---------------------------|-----------|----------|
| 1 | 9436 | Winnie the Pooh | Hardcover | English |
| 2 | 1029 | Le Petit Prince | Paperback | French |
| 3 | 8733 | Alice in Wonderland | Paperback | English |
| 4 | 2059 | Wind in the Willows | Hardcover | English |
| 5 | 5995 | A Bear Called Paddington | Paperback | English |
| 6 | 1031 | Pierre Lapin | Hardcover | French |
| 7 | 3984 | Le avventure di Pinocchio | Hardcover | Italian |

- ▶ Bitmap indices for Binding and Language have the following two and three position bitmaps:

Binding:

Hardcover: 1001011

Paperback: 0110100

Language:

English: 1011100

French: 0100010

Italian: 0000001

- ▶ A bitmap index makes it very fast to locate rows with a certain value.

Bitmap Index

- ▶ It is also possible to combine bitmap indexes
- ▶ Example: to find all hardcover books in English

Binding:

Hardcover: 1001011

Paperback: 0110100

Language:

English: 1011100

French: 0100010

Italian: 0000001

1001011

AND 1011100

= 1001000

- ▶ OR, NOT operations are also possible
- ▶ When there are few values, a bitmap index occupies little space
- ▶ It is also possible to make the bitmap index occupy less space: skip trailing zeros, various kinds of compression

- ▶ A **join index** is another means of speeding up query performance
- ▶ Like a materialized view, a join index represents the pre-computation of a query → **a join query**
- ▶ Unlike a materialized view, a join index contains pointers to the rows in the argument tables that satisfy the join predicate.
- ▶ A join index is capable of speeding up **certain queries** dramatically, while offering no speed up of other queries.
- ▶ Instead of storing list of pointers to rows, position bitmaps can be used.

Join Index

Example: Join Index

Book table

| RowID | BookID | Title | Genre |
|-------|--------|--------------------|------------------|
| 1 | 7493 | Tropical Food | Cooking |
| 2 | 9436 | Winnie the Pooh | Childrens' books |
| 3 | 9948 | Gone with the Wind | Fiction |
| 4 | 9967 | Italian Food | Cooking |

Sales table

| RowID | BookID | ShopID | SalesID | DayID | Count | Price |
|-------|--------|--------|---------|-------|-------|-------|
| 1 | 9436 | 854 | 1021 | 2475 | 2 | 30 |
| 2 | 7493 | 854 | 1021 | 2475 | 1 | 20 |
| 3 | 9948 | 876 | 2098 | 3456 | 1 | 20 |
| 4 | 7493 | 876 | 2231 | 3456 | 2 | 40 |
| 5 | 7493 | 876 | 3049 | 2475 | 1 | 20 |
| 6 | 9436 | 854 | 3362 | 3569 | 2 | 30 |
| 7 | 9967 | 731 | 3460 | 3569 | 1 | 35 |
| 8 | 7493 | 731 | 3460 | 3569 | 1 | 15 |
| 9 | 9948 | 731 | 3460 | 3569 | 1 | 15 |

Join index for Book and sales

list of pointers

| Book_RowID | Sales_RowID |
|------------|-------------|
| 1 | (2,4,5,8) |
| 2 | (1,6) |
| 3 | (3,9) |
| 4 | (7) |

position bitmaps

| Book_RowID | Sales_RowID |
|------------|-------------|
| 1 | 010110010 |
| 2 | 100001000 |
| 3 | 001000001 |
| 4 | 000000100 |

Outline

- 1 Hash Index
- 2 Bitmap and Join Indexes
- 3 Bloom Filters

Set membership queries

Given: a set M of n keys

- ▶ Each from some finite domain U

Sought: a data structure that supports **set membership queries**: is $k \in M$

- ▶ Common problem: caches, dictionaries, databases, ...
- ▶ Goal: make data structure as small as possible

Does this help? Not really, can show

- ▶ Essentially need to store M
- ▶ $\Omega(n \log(|U|/n))$ space

Approximate set membership

- ▶ Suppose we can tolerate a small amount of error
 - Few **false positives** acceptable
answer is yes, but $k \notin M$
 - No **false negatives** allowed
answer is no, but $k \in M$
- ▶ Now, we can do better: **Bloom Filters**
 - Bit vector of m bits, initialized 0
 - Hash function $h : U \mapsto \{0, \dots, m-1\}$
 - Insert key k : set bit $h(k)$
 - Query key k : answer yes, if bit $h(k)$ set, else no

Bloom Filters in practice

- ▶ Need to ensure low **false positive rate (FPR)**
 - Probability of false positive when asking for random $k \notin M$
 - use s hash functions instead of just one
 - Idea: Hash each key s times
 - Insert: set all s bits
 - Query: test whether all s bits = 1
 - Both operations take $O(s)$ time
- ▶ Assumptions
 - Storage cost of hash functions negligible
 - Hash values are random (independent, uniform)
 - FPR depends only on m, n, s
 - Theoretical analysis
some other time (look out for advanced data management courses next semester)

Evaluating an Index

- ▶ Which **access type** does the index support?
- ▶ Complexity of **accessing** a record
- ▶ Complexity of **inserting** a record
- ▶ Complexity of **deleting** a record
- ▶ **Space** complexity

Common Practice

- ▶ For point queries use hash index
- ▶ For point and range queries use B-tree
- ▶ Additionally, for multiple selections use Bitmaps and
- ▶ for set containment queries (with acceptable errors) use Bloom filters