

COL362/632 Introduction to Database Management Systems

Query Processing – Scans/Selections and Projections

Kaustubh Beedkar

Department of Computer Science and Engineering
Indian Institute of Technology Delhi




Outline

- 1 Query Processing
- 2 Selection and Scans
 - File Scan
 - Index Scan
- 3 General Selections
- 4 Selectivity Estimation
- 5 Projections


Query Processing (Overview)

- ▶ SQL : What?
- ▶ Physical plan: How?
- ▶ Query processing and optimization “compiles” a logical plan in to a physical plan

```
SELECT c.name, o.price, o.order_date  
FROM customer c, orders o  
WHERE c.customer_id = o.customer_id  
AND o.price > 50  
ORDER BY o.order_date DESC;
```




Logical Plan



Physical Plan

Query Processing (Overview)



Cost of an Operation

Cost depends on

- ▶ CPU cost depends on
 - # tuples
 - # index entries
 - # operators or functions
- ▶ I/O cost depends on
 - # blocks transferred by disk (b)
 - # random I/Os (S)

We usually ignore CPU cost (for disk-based systems)

- ▶ cost of an operation = $b \times t_T + S \times t_S$
- ▶ where, t_T is the time to transfer a block
- ▶ and t_S is the block access time (recall: seek time plus rotational delay)

Cost of an Operation

Assumptions and heuristics

- ▶ cost of accessing a random page = $\frac{1}{10}^{th}$ of the actual cost (why?)
- ▶ $B+$ Tree nodes are all in memory, and using a $B+$ Tree only incurs I/O cost for leaf node
- ▶ ...

Optimizer's goal is to minimize the **response time**

- ▶ Hard to estimate response time – depends on contents on the buffer, disks, data distribution
- ▶ Instead, minimize **resource consumption** (e.g., I/O cost)

- ▶ **Access path:** a “path” through which data can be located and accessed
- ▶ Scans are search algorithms that locate and retrieve “certain” records
 1. File scan
 2. Index scan

File Scan

Selection ($\sigma_{A=10}$)

Linear search

- ▶ $\text{cost} = t_S + b_r \times t_T$

Linear search (if A is key)

- ▶ $\text{cost} = t_S + \frac{b_r}{2} \times t_T$ (average case)

Search algorithm that uses an index

- ▶ Use an index that is available on some predicate
- ▶ Cost depends on the index

Selection ($\sigma_{A=10}$)

Cost depends on height plus type of index

1. Clustering Index

Case: A is key

► $\text{cost} = (h_i + 1) \times (t_T + t_S)$

Case: A is not a key

► $\text{cost} = h_i \times (t_T + t_S) + t_S + (b \times t_T)$

Selection ($\sigma_{A=10}$)

Cost depends on height plus type of index

2. Non-clustering Index

Case: A is key

► $\text{cost} = (h_i + 1) \times (t_T + t_S)$

Case: A is not a key

► $\text{cost} = (h_i + n) \times (t_T + t_S)$

n = number of records fetched

Index Scan – B+Tree

Selection ($\sigma_{A>10}$ or $\sigma_{A\geq 10}$)

1. Clustering Index

- ▶ $\text{cost} = h_i \times (t_T + t_S) + t_S + b \times t_T$
- ▶ Note: for $\sigma_{A<v}$ or $\sigma_{A\leq v}$, an index is not useful (Why?)

2. Non-clustering Index

- ▶ $\text{cost} = (h_i + n) \times (t_T + t_S)$
- ▶ Only use when n is small
(HW: compare with Linear search. What would happen if n is large?)
- ▶ Often used with Bitmaps and sorting (HW: Read about Bitmap Index Scan in PostgreSQL)

General Selections

- ▶ So far we only looked at σ_p where $p \equiv A \theta v$
- ▶ Predicates often involve
 1. conjunctions, i.e., $\sigma_{p_1 \wedge p_2 \wedge \dots \wedge p_n}$ e.g., $R.A = 10 \text{ AND } R.B < 20$
 2. disjunctions, i.e., $\sigma_{p_1 \vee p_2 \vee \dots \vee p_n}$ e.g., $R.A > 10 \text{ OR } R.B < 20$
- ▶ General selections require **Index Matching**
 - an index matches a selection predicate if the index can evaluate it
 - e.g, $R(A,B,C)$ and a hash index on composite key (A,B)
 $\sigma_{A=5 \wedge B=10}$ matches the index
 $\sigma_{A=5}$ **does not** matches the index

Index Matching – Hash Index

Selection ($\sigma_{p_1 \wedge p_2 \wedge \dots \wedge p_n}$)

A hash index on composite key (A,B,..) matches the selection condition if

- ▶ **all** attributes in the index search key appear in the predicate with $\theta \in \{=\}$

Example R(A,B,C)

selection condition	hash index on (A,B)	hash index on (A)
A = 5 and B = 10	yes	yes
A = 5	no	yes
C = 5	no	no
A > 5 and B < 10	no	no
A = 5 and B = 10	yes	yes
A = 5 and B = 10 and C < 1	yes	yes

Index Matching – B+Tree

Selection ($\sigma_{p_1 \wedge p_2 \wedge \dots \wedge p_n}$)

A B+Tree index on composite key (A,B,..) matches the selection condition if

- ▶ the attributes in the predicate form a prefix of the search key of B+Tree
- ▶ $\theta \in \{=, <, >, \leq, \geq\}$

Example R(A,B,C)

selection condition	B+Tree index on (A,B)	B+Tree index on (B,C)
A = 5 and B = 10	yes	yes
A = 5	yes	no
B = 5	no	yes
A > 5 and B < 10	yes	yes
A = 5 and B = 10	yes	yes
A = 5 and B = 10 and C < 1	yes	yes

Index Matching

When more than one index can match

Example

- ▶ B+Tree index on (A,B), and a hash index on (C)
- ▶ Selection condition: $A = 5$ and $B = 3$ and $C = 10$

Which index to use?

[Option 1] Use B+Tree index, retrieve tuples with ($A=5$ and $B=3$), then check $C=10$ for every retrieved tuple

[Option 2] Use hash index, retrieve tuples with ($C=10$), then check $A=5$ and $B=3$ for every retrieved tuple

[Option 3] Use B+Tree index, retrieve tuples with ($A=5$ and $B=3$), use hash index and retrieve tuples with $C=10$, intersect the record ids, only then fetch the tuples from disk

Disjunctive selection by union of identifiers

Example

- ▶ Hash index on (A) and B+Tree index on (B)
- ▶ Selection condition: $A = 5$ or $B > 10$
- ▶ [Option 1] Do a linear scan
- ▶ [Option 2] Use both indexes, fetch the records ids, and do a union, only then retrieve the tuples

Example

- ▶ Hash index on (A) + hash index on (B)
- ▶ Selection condition: $A = 5$ or $B > 10$
- ▶ [Option 1] Only option is to do a linear scan

Selectively Estimation

- ▶ Selectivity = fraction of tuples that need to be retrieved
- ▶ Goal should to chose the **most selective path** (Why?)
- ▶ But, estimating selectivity of an access path is a hard (active research area!)

Selectivity Estimation

Example

- ▶ hash index on (A, B)
- ▶ selection condition: A=3 and B=6

$$\text{Selectivity} \approx \frac{1}{\#keys}$$

- ▶ Number of keys can be determined from the index
- ▶ Only holds under uniform distribution assumption
- ▶ selection condition: A=3 and B=6 and C=10
- ▶ If no knowledge of $\#keys$
 - selectivity = product of selectivity of primary conjuncts
 - or, use $\frac{1}{10}$ as default
 - holds only under independence assumption

Selectivity Estimation

Example

- ▶ Selection condition: $25 < A < 50$

$$\text{selectivity} \approx \frac{\text{interval}}{\text{max} - \text{min}}$$

Simple Projections

```
select R.A, R.B, from ... where ...
```

- ▶ Scan the file for each tuple output R.A, R.B

Projection w/ Duplicate Elimination

`select distinct R.A, R.B from ... where ...`

- ▶ first do simple projections
- ▶ remove duplicates
 - Sort-based duplicate elimination
 - Optimization: eliminate duplicates when merging runs
 - Hash-based duplicate elimination