

COL362/632 Introduction to Database Management Systems


Query Processing – Joins

Kaustubh Beedkar

Department of Computer Science and Engineering
Indian Institute of Technology Delhi



Query Processing (Overview)



Outline

- 1 Join Basics
- 2 Nested Loop Join
- 3 Sort-Merge Join
- 4 Hash Join

Joins


$$R \bowtie_{\theta} S$$

- ▶ We will assume equijoins, i.e., $\theta \equiv \{ R.A = S.A \}$
- ▶ Equijoin algorithms can be extended to support other joins
- ▶ R is **outer relation**
- ▶ S is **inner relation**

$R \bowtie S$ is one of the most common operation and needs to be carefully optimized

- ▶ Many algorithms
- ▶ No one size fits all

```
SELECT c.name, o.price, o.order_date
FROM customer c, orders o
WHERE c.customer_id = o.customer_id
AND o.price > 50
ORDER BY o.order_date DESC;
```



Logical Plan

Join Algorithms


1. Nested Loop Join (NLJ)

- Naive NLJ
- Block NLJ
- Index NLJ

2. Sort-Merge Join

3. Hash Join

- Simple hash join
- Grace hash join
- Hybrid hash join



Notations & Cost Factors

► For relation R

- r a tuple in R
- b_R blocks
- n_R tuples

► For Relation S

- s a tuple in S
- b_S blocks
- n_S tuples

customer		
c_id	name	address
3	A	Delhi
2	B	Delhi
1	C	Mumbai
5	D	Patna
4	E	Surat
6	F	Puri
7	G	Mumbai

order			
o_id	c_id	price	Date
o1	2	100	2020
o2	4	40	2021
o3	2	30	2023
o4	1	75	2024
o5	3	50	2022

► Cost model based on # of I/Os

- #blocks trasferred (recall: each block requires t_T time)
- #blocks seeked (recall: each block requires t_S time)

Naive Nested Loop Join

Algorithm

```
1: for  $r \in R$  do  
2:   for  $s \in S$  do  
3:     if if  $r$  and  $s$  match then  
4:        $\text{EMIT}(r \bowtie s)$ 
```

customer		
c_id	name	address
3	A	Delhi
2	B	Delhi
1	C	Mumbai
5	D	Patna
4	E	Surat
6	F	Puri
7	G	Mumbai

order			
o_id	c_id	price	Date
o1	2	100	2020
o2	4	40	2021
o3	2	30	2023
o4	1	75	2024
o5	3	50	2022

Cost

- ▶ #blocks transferred = $b_R + (n_R \times b_S)$
- ▶ #seeks = $b_R + n_R$
- ▶ cost = $[b_R + (n_R \times b_S)] \cdot t_T + (b_R + n_R) \cdot t_S$

Naive Nested Loop Join

Example

- ▶ $R : b_R = 100; n_R = 5000$
- ▶ $S : b_S = 400; n_S = 10,000$

Cost

- ▶ $b_R + (n_R \times b_S) = 100 + (5000 \times 400) = 2,000,100$
- ▶ $b_R + n_R = 100 + 5000 = 5010$
- ▶ At 4 ms/block seek time and 0.1 ms/block transfer time, total time $\approx 3.5\text{min}$

If we consider S as the outer relation

- ▶ $b_S + (n_S \times b_R) = 400 + (10,000 \times 100) = 1,000,400$
- ▶ $b_S + n_S = 400 + 10,000 = 10,400$
- ▶ At 4 ms/block seek time and 0.1 ms/block transfer time, total time $\approx 2.3\text{min}$

Block Nested Loop Join

Algorithm

```
1: for each block  $B_R \in R$  do
2:   for each block  $B_S \in S$  do
3:     for each  $r \in B_R$  do
4:       for each  $s \in B_S$  do
5:         if  $r$  and  $s$  match then
6:            $\text{EMIT}(r \bowtie s)$ 
```

customer		
c_id	name	address
3	A	Delhi
2	B	Delhi
1	C	Mumbai
5	D	Patna
4	E	Surat
6	F	Puri
7	G	Mumbai

order			
o_id	c_id	price	Date
o1	2	100	2020
o2	4	40	2021
o3	2	30	2023
o4	1	75	2024
o5	3	50	2022

Cost

- ▶ block transfers: $b_R + (b_R \times b_S)$
- ▶ seeks: $2b_R$

Block Nested Loop Join

- ▶ Reduces to access to disk
- ▶ Smaller relation should be the outer table (why?)
 - Smaller based on number of pages!

Optimizations

- ▶ If join attribute of inner relation is a key, terminate outer loop as soon as a match is found
- ▶ **Cache-conscious inner loop**: scan alternatively forward and backward (how does this help?)
- ▶ Use biggest size of a relation as blocking unit (next slide)

Block Nested Loop Join

Assume that we have B buffer pages


- ▶ Use $B - 2$ buffer pages for R
- ▶ Use one buffer page for S
- ▶ Use one buffer page for output

Algorithm

- 1: **for** each $B - 2$ blocks $B'_R \in R$ **do**
 - 2: **for** each block $B_S \in S$ **do**
 - 3: **for** each $r \in B - 2$ blocks **do**
 - 4: **for** each $s \in B_S$ **do**
 - 5: **if** r and s match **then**
 - 6: $\text{EMIT}(r \bowtie s)$
- ▶ Improved cost = $b_R + \left(\left\lceil \frac{b_R}{B-2} \right\rceil \times b_S \right)$
 - ▶ Improved seeks = $2 \left\lceil \frac{b_R}{B-2} \right\rceil$

customer		
c_id	name	address
3	A	Delhi
2	B	Delhi
1	C	Mumbai
5	D	Patna
4	E	Surat
6	F	Puri
7	G	Mumbai

order			
o_id	c_id	price	Date
o1	2	100	2020
o2	4	40	2021
o3	2	30	2023
o4	1	75	2024
o5	3	50	2022



Buffer pool ($B=4$)

and so on...

Block Nested Loop Join

Example

- ▶ $R : b_R = 100; n_R = 5000$
- ▶ $S : b_S = 400; n_S = 10,000$

Cost

If $B = 12$

- ▶ $b_R + \left(\left\lceil \frac{b_R}{B-2} \right\rceil \times b_S \right) = 100 + 10 \times 400 = 4,100$
- ▶ $2 \left\lceil \frac{b_R}{B-2} \right\rceil = 2 \times 10 = 20$
- ▶ At 0.1ms/block transfer time and 4ms/block seek time, total time $\approx 0.4s$

If $b_R < B - 2$ (outer relation will fit in memory)

- ▶ 500 block transfers
- ▶ 2 seeks
- ▶ At 0.1ms/block transfer time and 4ms/block seek time, total time $\approx 0.05s$

Index Nested Loop Join

Use an index scan instead of sequential scan for inner relation

Algorithm

- 1: **for** each $r \in R$ **do**
- 2: **for** each $s \in \text{INDEX}(r_i = s_j)$ **do**
- 3: **if** r and s match **then**
- 4: $\text{EMIT}(r \bowtie s)$

Cost

- ▶ b_R block transfers + b_R seeks
- ▶ index scan cost: $n_R \times C$
- ▶ $\text{cost} = b_R(t_T + t_S) + n_R \times C$
 - Assuming C is the cost of each index probe (recall cost of single selection)
- ▶ If index on both relations, use one with fewer tuples as outer relation

customer		
c_id	name	address
3	A	Delhi
2	B	Delhi
1	C	Mumbai
5	D	Patna
4	E	Surat
6	F	Puri
7	G	Mumbai

order			
o_id	c_id	price	Date
o1	2	100	2020
o2	4	40	2021
o3	2	30	2023
o4	1	75	2024
o5	3	50	2022

Index
on c_id

Sort Merge Join

1. Sorting

- ▶ Sort relations on join key(s)

2. Merge

- ▶ Scan two sorted relations and emit matching tuples
 - Note: This is different than merging of external sort

Sort Merge Join

Algorithm

- 1: $R' \leftarrow \text{SORT}(R)$, $S' \leftarrow \text{SORT}(S)$ on join keys
- 2: $p_R \leftarrow$ address of first tuple in R'
- 3: $p_S \leftarrow$ address of first tuple in S'
- 4: **while** p_R and p_S **do**
- 5: **if** $p_R > p_S$ **then**
- 6: $p_S \leftarrow$ next tuple of S'
- 7: **if** $p_R < p_S$ **then**
- 8: $p_R \leftarrow$ next tuple of R'
- 9: backup if required
- 10: **if** p_R and p_S match **then**
- 11: $\text{EMIT}(r \bowtie s)$
- 12: $p_S \leftarrow$ next tuple of S'



customer		
c_id	name	address
1	C	Mumbai
2	B	Delhi
3	A	Delhi
4	E	Surat
5	D	Patna
6	F	Puri
7	G	Mumbai



order			
o_id	c_id	price	Date
o4	1	75	2024
o1	2	100	2020
o3	2	30	2023
o5	3	50	2022
o2	4	40	2021

EOF

Sort Merge Join

Cost

- ▶ Sorting R : $2b_R(1 + \lceil \log_{B-1} \lceil \frac{b_R}{B} \rceil \rceil)$
- ▶ Sorting S : $2b_S(1 + \lceil \log_{B-1} \lceil \frac{b_S}{B} \rceil \rceil)$
- ▶ Merging: $b_R + b_S$ (block transfers)
- ▶ Seeking: $\lceil \frac{b_R}{b_B} \rceil + \lceil \frac{b_S}{b_B} \rceil$
 - b_B buffer blocks allocated to each relation

Sort Merge Join

Example

- ▶ $R : b_R = 100; n_R = 5000$
- ▶ $S : b_S = 400; n_S = 10,000$

Cost

- ▶ if $B=11$ and $b_B = 1$
- ▶ sorting cost for $R : 2 \times 100(1 + \lceil \log_{10} \lceil 100/11 \rceil \rceil) = 400$
- ▶ sorting cost for $S : 2 \times 400(1 + \lceil \log_{10} \lceil 400/11 \rceil \rceil) = 2,400$
- ▶ merging cost = 500
- ▶ seeking: cost = 500
- ▶ At 0.1ms/transfer seek time and 4ms/block seek time, total time $\approx 2.3s$

Sort Merge Join

Notes

- ▶ Worst case: when join attribute of all tuples contains the same value
- ▶ Useful when table(s) is already sorted in join key
- ▶ Useful when sorted output is required

Basic Idea

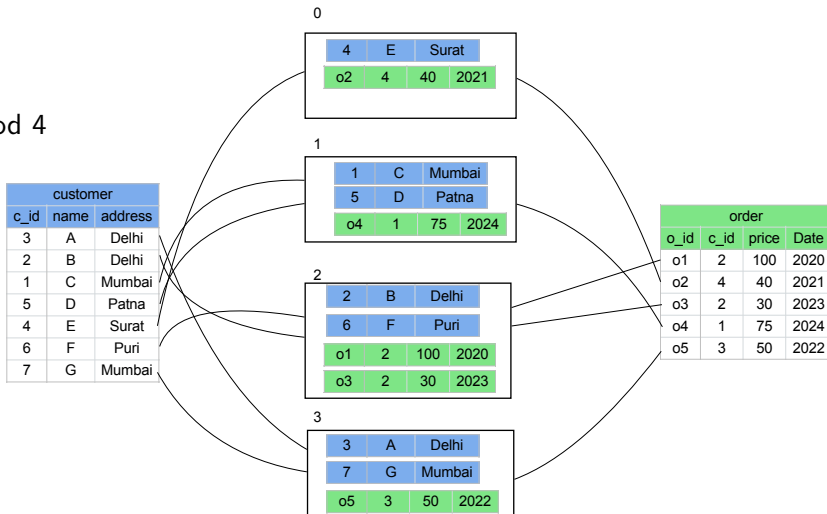
Key insight

- ▶ if $r \in R$ matches $s \in S$, then both tuples when hashed on the join attribute using the same hash function should be in the same bucket i
- ▶ Partition R into R_0, R_1, \dots, R_n
- ▶ Partition S into S_0, S_1, \dots, S_n
- ▶ Using a hash function $h : \text{join_attribute} \mapsto \{0, 1, \dots, n\}$

Basic Idea

Example

► $h : x \bmod 4$



Naive In-memory Hash Join

1. **Build:** Scan the outer relation and build a hash table
2. **Probe:** Scan the inner relation and find the matching tuple

Algorithm

- 1: build a hash table H for R
- 2: **for** each $s \in S$ **do**
- 3: **if** $h(s) \in H$ **then**
- 4: find matching tuple(s) and **EMIT**


- ▶ Does not work on large relations
- ▶ Buffer pool manager may swap out pages of hash table!

Grace Hash Join

- ▶ Also known as partitioned hash join
- ▶ Adopts a divide and conquer approach
- ▶ **Partitioning phase:** Partition R and S using the same hash function
- ▶ **Build & Probe Phase:** Compare tuples in each partition


GRACE Parallel Relational Database Machine

<https://museum.ipsj.or.jp/en/computer/other/0014.html>



Grace Hash Join (Example)

- Compute $R \bowtie S$



1 Buffer page to read





B-1 Buffer pages for partitioning



Grace Hash Join (Example)– Partitioning Phase

► Partition S_R



1 Buffer page to read



B-1 Buffer pages for partitioning






1 Buffer page to read



Grace Hash Join (Example)– Partitioning Phase

► Partition R



1 Buffer page to read



B-1 Buffer pages for partitioning




1 Buffer page to read



Grace Hash Join (Example)

- After partitioning R and S , we have







1 Buffer page to read



B-1 Buffer pages for partitioning

Grace Hash Join (Example) – Build & Probe Phase


- Build in-memory hash table on R



1 Buffer page to read






B-2 Buffer pages hash table



Grace Hash Join (Example) – Build & Probe Phase

- Build in-memory hash table on R



1 Buffer page to read



B-2 Buffer pages hash table



1 Buffer page for output

1 Buffer page to read





B-2 Buffer pages hash table



Grace Hash Join (Example) – Build & Probe Phase


- Probe S and compute matching tuples



1 Buffer page to read



B-2 Buffer pages hash table



1 Buffer page for output

1 Buffer page to read






B-2 Buffer pages hash table



Grace Hash Join (Example) – Build & Probe Phase

- ▶ Same process for other partitions



1 Buffer page to read



B-2 Buffer pages hash table



1 Buffer page for output

1 Buffer page to read

Corner Cases

If $n \geq \#$ blocks of memory

- ▶ Relations cannot be partitioned in one pass
- ▶ in one pass, a relation can at most be partitioned into $B - 1$ partitions (or number of buffers available as output buffers)
- ▶ if a partition does not fit in memory, then **recursively partition** using different hash functions

If partitioning is **skewed**

- ▶ Increase the number of partitions by a **fudge factor** (usually 20%)
- ▶ **Overflow resolution** by re-partitioning
- ▶ **Overflow avoidance** by first create many smaller partitions, then combine them to fit in memory
- ▶ if resolution and avoidance fail, then use NLJ for keys with “many” values

Partitioned Hash Join

Cost (if no recursive partitioning is required)

- ▶ partitioning phase : $2 \times (b_R + b_S)$ block transfers
- ▶ build & probe phase : $b_R + b_S + 4n$
 - $4n$ is the overhead for partially filled blocks (usually ignored in cost calculations)
- ▶ Total transfer cost = $3(b_R + b_S)$

- ▶ $2(\lceil \frac{b_R}{b_B} \rceil + \lceil \frac{b_S}{b_B} \rceil)$ seeks in partitioning phase
- ▶ $2n$ seeks in build and probe phase
- ▶ Total seek cost = $2(\lceil \frac{b_R}{b_B} \rceil + \lceil \frac{b_S}{b_B} \rceil) + 2n$

Partitioned Hash Join

Example

- ▶ $R : b_R = 100; n_R = 5000$
- ▶ $S : b_S = 400; n_S = 10,000$

Cost

- ▶ assuming $b_B = 3$
- ▶ $3(100 + 400) = 1500$ block transfers
- ▶ $2(\lceil \frac{100}{3} \rceil + \lceil \frac{400}{3} \rceil) = 336$ seeks
- ▶ At 0.1ms/block transfer time and 4ms/block seek time, total time ≈ 0.6 s

Note

- ▶ If memory is large enough, simple hash join requires only $b_R + b_S$ block transfers and 2 seeks!

Hash Join

Notes

- ▶ We don't care about size of inner table, only outer table needs to fit in memory
- ▶ Use static hashing, if size of outer table is known
otherwise, use dynamic hashing (Recall extendible hashing)

Homework exercise

Join Algorithm	I/O cost (block transferred)	Example
NLJ		
BNLJ		
SMJ		
GHJ		

Other Operators

Self study: aggregation and set operators (B1: Ch 15 (15.6))