

# COL362/632 Introduction to Database Management Systems

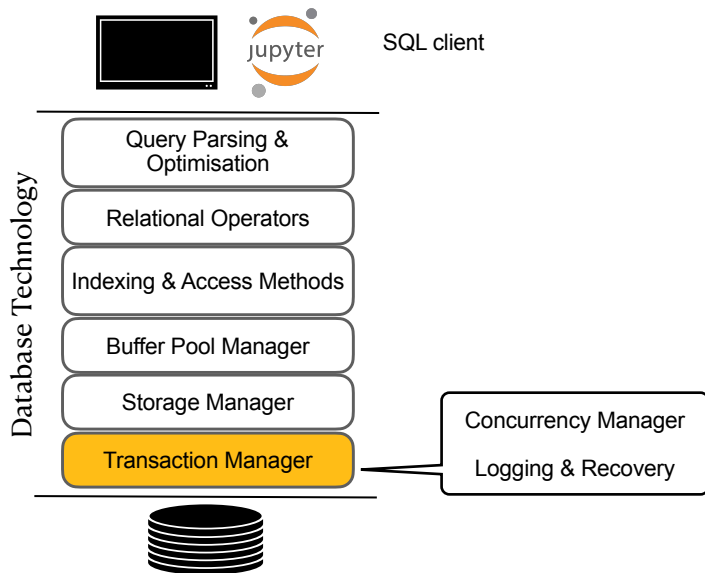
## Recovery

Kaustubh Beedkar

Department of Computer Science and Engineering  
Indian Institute of Technology Delhi




# Introduction



# Recall Buffer Pool Manager


- ▶ We use Buffer Pool Manager to read and write to disk
- ▶ Actual writing to disk is dictated by the buffer replacement policy



# Failure Scenario I

- ▶ Consider the following schedule

$T_1$   
—  
read(A)  
A:=A-50  
write(A)  
...  
commit




- ▶ Transaction committed before page was flushed to disk—(Durability challenges!)

## Failure Scenario II

- Consider the following schedule

$T_1$	$T_2$
read(A)	
write(A)	
...	read(B)
	write(B)
	commit
abort	



- Changes of uncommitted Tx flushed to disk—(Atomicity challenges!)

# Why do Transactions Fail?

## 1. Transaction Failures

- Logical Errors (e.g., some integrity constraint violation)
- Internal State Errors (e.g., deadlock)

## 2. System Failure

- Software failures – OS or DBMS implementation (e.g., uncaught divide by zero)
- Hardware failures (e.g., power goes off)
- **Fail-stop assumption** – non-volatile storage contents are assumed to not be corrupted by system crash

## 3. Storage Media Failures

- Disk failure (e.g., a disk head crash)
  - destruction can be detected (use checksums to detect failures)
- Note: the database cannot recover from this! Restore from an archived version

## DBMS must ensure

**ACID:** Changes from any Tx are durable once it has been committed

**ACID:** No partial changes are durable if the Tx are aborted

## Logging & recovery

- ▶ Actions taken during normal execution to ensure that DBMS can recover from failure
- ▶ Actions taken after a failure to recover the database to a state that ensures consistency, atomicity, and durability

# Redo & Undo

The two issues

Changes from any Tx are durable once it has been committed

- ▶ **Redo** certain actions that were committed but not written to disk

No partial changes are durable if the Tx are aborted

- ▶ **Undo** certain actions that did not commit



# Storage Types

## ► Volatile Storage

- Does not survive system crash
- e.g., DRAM, cache memory


## ► Non-volatile Storage

- Survives system crash
- e.g., HDD, SSD

## ► Stable Storage

- A mythical form of storage that survives all possible failures
- Approximated by replication

## Storage hierarchy



# Some Terminology

- ▶  $\text{read}(A)$
- ▶  $\text{write}(A)$
- ▶  $\text{input}(B_A)$ 
  - Transfer the physical block  $B_A$  where the data items  $A$  resides to buffer pool
- ▶  $\text{output}(B_A)$ 
  - Transfer the buffer pool page  $B_A$  that had the data item  $A$  to disk (replace the old block  $B_A$  on disk)
  - Note:  $\text{output}(B_A)$  may not happen immediately following  $\text{write}(A)$

## In the context of transactions

- ▶ A Tx must perform  $\text{read}(A)$  before accessing  $A$  for the first time
- ▶  $\text{write}(A)$  can be executed anytime before the transaction commits

# Log-based Recovery

## Key Idea

Output first the information describing the modifications to stable storage without modifying the database

## Log File

- ▶ A log file contains the changes
- ▶ When a Tx  $T_i$  starts, it registers itself by writing a  $\langle T_i, start \rangle$  log record
- ▶ Before  $T_i$  executes a write(A), a log record  $\langle T_i, A, V_{old}, V_{new} \rangle$  is written
- ▶ When  $T_i$  finishes its last statement, the log record  $\langle T_i, commit \rangle$  is written

Logging allows to perform both undo and redo operations

## **Immediate-modifications**

- ▶ Perform updates to buffer pool/disk before the Tx commits

## **Deferred-modifications**

- ▶ Perform updates to buffer pool/disk only at the time of Tx commit

Note: All modification must be preceded by creation of a log record

# Concurrency & Recovery

- ▶ Consider this schedule

$T_1$	$T_2$
...	
write(A)	
	...
	write(A)
abort	
	abort

- ▶ If  $T_1$  aborts, we undo( $T_1$ ), and also have to undo( $T_2$ )
- ▶ Require that if data item has been modified by a transaction, no other transaction can modify the data item until the first transaction commits or aborts
  - automatically happens in 2PL
  - time-stamp ordered protocol (recall dealing with cascaded aborts)
  - validation-based also supports this

## When does a Tx go in committed state?

- ▶ when the commit log record  $\langle T_i, commit \rangle$  (last record of a Tx) has been written to stable storage
- ▶ we don't care if the database has been modified or not

Also, recall partially committed state

# Undo and Redo Operation

## **redo( $T_i$ )**

- ▶ Sets the value of all data items updated by  $T_i$  to the new values, going from first log record for  $T_i$
- ▶ No logging is done in this case

## **undo( $T_i$ )**

- ▶ Restore the value of all data items updated by  $T_i$  to their old value, going backwards from the last record of  $T_i$ 
  - Also write a **redo-only** record  $\langle T_i, A, V_{old} \rangle$
  - When an undo( $T_i$ ) finishes, a  $\langle T_i, abort \rangle$  log record is written out

## When recovering from failure

- ▶  $T_i$  needs to be undone if the log
  - contains the record  $\langle T_i, start \rangle$
  - but does not contain either the record  $\langle T_i, commit \rangle$  or  $\langle T_i, abort \rangle$
- ▶  $T_i$  needs to be redone if the log
  - contains the records  $\langle T_i, start \rangle$
  - and either the record  $\langle T_i, commit \rangle$  or  $\langle T_i, abort \rangle$



# Example

Consider the schedule

$T_1$	$T_2$
begin	
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
	begin
	read(C)
	$C := C - 100$
	write(C)

Scenario-1: System crash after write(B)

Log records

$\langle T_1 \text{start} \rangle$

$\langle T_1, A, 100, 50 \rangle$

$\langle T_1, B, 100, 150 \rangle$

Recovery action:

- ▶ undo( $T_1$ ): B is restored to 100; A to 100
- ▶ and write log records  
 $\langle T_1, B, 100 \rangle, \langle T_1, A, 100 \rangle, \langle T_1, \text{abort} \rangle$

# Example

Consider the schedule

$T_1$	$T_2$
begin	
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
commit	
	begin
	read(C)
	$C := C - 100$
	write(C)

Scenario-2: System crash after write(C)

Log records

$\langle T_1, start \rangle$   
 $\langle T_1, A, 100, 50 \rangle$   
 $\langle T_1, B, 100, 150 \rangle$   
 $\langle T_1, commit \rangle$   
 $\langle T_2, start \rangle$   
 $\langle T_2, C, 150, 50 \rangle$

Recovery action:

- ▶ redo( $T_1$ ) and undo( $T_2$ ): A is set to 50; B to 150, and C is restored to 150
- ▶ and write log records  
 $\langle T_2, C, 150 \rangle, \langle T_2, abort \rangle$

# Example

Consider the schedule

$T_1$	$T_2$
begin	
read(A)	
$A := A - 50$	
write(A)	
read(B)	
$B := B + 50$	
write(B)	
commit	
	begin
	read(C)
	$C := C - 100$
	write(C)
	commit

Scenario-3: System crash after  $T_2$  commits

Log records

$\langle T_1, start \rangle$   
 $\langle T_1, A, 100, 50 \rangle$   
 $\langle T_1, B, 100, 150 \rangle$   
 $\langle T_1, commit \rangle$   
 $\langle T_2, start \rangle$   
 $\langle T_2, C, 150, 50 \rangle$   
 $\langle T_2, commit \rangle$

Recovery action:

- ▶ redo( $T_1$ ) and redo( $T_2$ ): A is set to 50; B to 150. Then C is set to 50

# Checkpoints

- ▶ Redo and Undo operations can have significant overhead
  - Processing entire log can be time consuming
  - Redoing Tx that have already written to disk is wasteful
- ▶ **Checkpoints**
  - Special marker: recovery only needs to look at parts of the log after checkpoint and just before the checkpoint
- ▶ Checkpointing is done periodically
  - output all log records from buffer pool to stable storage
  - output all modified buffer blocks to disk
  - Write a log record  $\langle \textit{checkpoint}, L \rangle$ , where  $L$  is a list of all transactions active at the time of checkpoint

No Tx allowed to update while checkpointing is in progress

# Write-Ahead Logging

## Key Ideas

- ▶ Tx  $T_i$  enters commit state only after  $\langle T_i, commit \rangle$  log record has been output to disk
- ▶ Before the  $\langle T_i, commit \rangle$  log record can be output to stable storage, all log records must have been output to stable storage
- ▶ Write to disk the log file records corresponding to database modifications **before** the buffer pool manager flushes pages to disk
- ▶ Buffer Pool Policy: Steal + No-force
  - Steal policy: Pages of uncommitted Tx may be written to disk
  - No-force policy: Allow Tx to commit even if it has modified pages that have not been written to disk

# Write-Ahead Logging

Consider the schedule

$T_1$

begin

read(A)

write(A)


read(B)

write(B)

...

commit


- ▶ Assume that the modified page (with  $A = 50$ ) had to be replaced
- ▶ First write the log page to disk, then flush the page




# Write-Ahead Logging (scenario 2)

Consider the schedule

$T_1$   
begin  
read(A)  
write(A)  
read(B)  
write(B)  
...  
commit



- ▶ Assume that  $T_1$  is ready to commit
- ▶ First write the log page to disk, then commit the  $T_x$



## Optimization

- ▶ Group commits: batch multiple commits

# Recovery Algorithm (sketch)

## **Write ahead logging** (during normal execution)

- ▶  $\langle T_i \text{ start} \rangle$  when the Tx begins
- ▶  $\langle T_i, A, V_{old}, V_{new} \rangle$  for each write
- ▶  $\langle T_i \text{ commit} \rangle$  when Tx ends



# Recovery Algorithm (sketch)

## In case of transaction failure

assuming  $T_i$  needs to be rolled back

- ▶ Scan log backwards
- ▶ For each log record of the form  $\langle T_i, X, V_{old}, V_{new} \rangle$ 
  - perform undo by writing  $V_{old}$  to  $X$
  - write a **compensation log record**  $\langle T_i, X, V_{old} \rangle$
- ▶ After  $\langle T_i \text{ start} \rangle$  is found, stop scanning and write log record  $\langle T_i \text{ abort} \rangle$

# Recovery Algorithm (sketch)

## In case of system failure

1. Redo phase: replay updates of **all** transactions

- ▶ Find last checkpoint  $\langle \text{checkpoint } L \rangle$  record and set undo-list to  $L$
- ▶ Scan log forward from  $\langle \text{checkpoint } L \rangle$  record
  - whenever a record  $\langle T_i, X, V_{old}, V_{new} \rangle$  is found, redo by writing  $V_{new}$  to  $X$
  - whenever a record  $\langle T_i \text{ start} \rangle$  is found, add  $T_i$  to undo-list
  - whenever a record  $\langle T_i \text{ commit} \rangle$  or  $\langle T_i \text{ abort} \rangle$  is found, remove  $T_i$  from undo-list

# Recovery Algorithm (sketch)

## In case of system failure

### 2. Undo phase:

- ▶ Scan log backwards from end
  - whenever a record  $\langle T_i, X, V_{old}, V_{new} \rangle$  is found, where  $T_i$  is in undo list
    - perform undo by writing  $V_{old}$  to  $X$
    - write a **compensation log record**  $\langle T_i, X, V_{old} \rangle$
  - whenever a record  $\langle T_i \text{ start} \rangle$  is found where  $T_i$  is in undo list
    - write log record  $\langle T_i \text{ abort} \rangle$
    - remove  $T_i$  from undo-list
  - stop when undo-list is empty

# Recovery Example

- ▶ Consider the following log

$\langle T_1 \text{ start} \rangle$   
 $\langle T_1, B, 100, 150 \rangle$   
 $\langle T_2 \text{ start} \rangle$   
 $\langle \text{checkpoint } \{ T_1, T_2 \} \rangle$   
 $\langle T_2, C, 100, 200 \rangle$   
 $\langle T_2 \text{ commit} \rangle$   
 $\langle T_3 \text{ start} \rangle$   
 $\langle T_3, A, 100, 50 \rangle$   
 $\langle T_1, B, 100 \rangle$   
 $\langle T_1 \text{ abort} \rangle$   
 $\langle T_3, A, 100 \rangle$   
 $\langle T_3 \text{ abort} \rangle$

Scan forward from checkpoint (redo phase)

- ▶ Undo-list:  $T_1$  ~~$T_1$~~ ,  $T_2$  ~~$T_2$~~ ,  $T_3$
- ▶ After  $\langle T_2 \text{ commit} \rangle$  Remove  $T_2$  from undo-list
- ▶ After  $\langle T_3 \text{ start} \rangle$  Add  $T_3$  to undo-list
- ▶ After  $\langle T_1 \text{ abort} \rangle$  Remove  $T_1$  from undo-list

Scan backward (undo phase)

- ▶ After writing 100 to A, Write compensation log record  $\langle T_3, A, 100 \rangle$
- ▶ After  $\langle T_3 \text{ start} \rangle$ , write  $\langle T_3 \text{ abort} \rangle$